

# Maximum Likelihood Estimation and Ordinary Least Squares on Simple Linear Regression

August 6, 2019

```
[3]: #!pip install --upgrade pip --user  
#!pip install numdifftools --user  
#!pip install scipy --upgrade --user  
  
!pip uninstall statsmodels -y --user  
!pip install statsmodels==0.10.0rc2 --pre --user
```

Usage:

```
pip uninstall [options] <package> ...  
pip uninstall [options] -r <requirements file> ...
```

no such option: --user

Collecting statsmodels==0.10.0rc2

Downloading [https://files.pythonhosted.org/packages/ab/f5/6bb191bb31574f59b42ee10278a8d002e5be4055d41f9af86d682b8c63b6/statsmodels-0.10.0rc2-cp36-cp36m-manylinux1\\_x86\\_64.whl](https://files.pythonhosted.org/packages/ab/f5/6bb191bb31574f59b42ee10278a8d002e5be4055d41f9af86d682b8c63b6/statsmodels-0.10.0rc2-cp36-cp36m-manylinux1_x86_64.whl) (8.1MB)

|| 8.1MB 8.3MB/s eta 0:00:01

| | 3.9MB 2.6MB/s eta 0:00:02

Requirement already satisfied: pandas>=0.19 in  
/home/data/.local/lib/python3.6/site-packages (from statsmodels==0.10.0rc2)  
(0.24.2)

Requirement already satisfied: patsy>=0.4.0 in  
/home/data/.local/lib/python3.6/site-packages (from statsmodels==0.10.0rc2)  
(0.5.1)

Requirement already satisfied: scipy>=0.18 in  
/home/data/.local/lib/python3.6/site-packages (from statsmodels==0.10.0rc2)  
(1.3.0)

Requirement already satisfied: numpy>=1.11 in  
/home/data/.local/lib/python3.6/site-packages (from statsmodels==0.10.0rc2)  
(1.17.0)

Requirement already satisfied: python-dateutil>=2.5.0 in /usr/lib/python3/dist-packages (from pandas>=0.19->statsmodels==0.10.0rc2) (2.6.1)

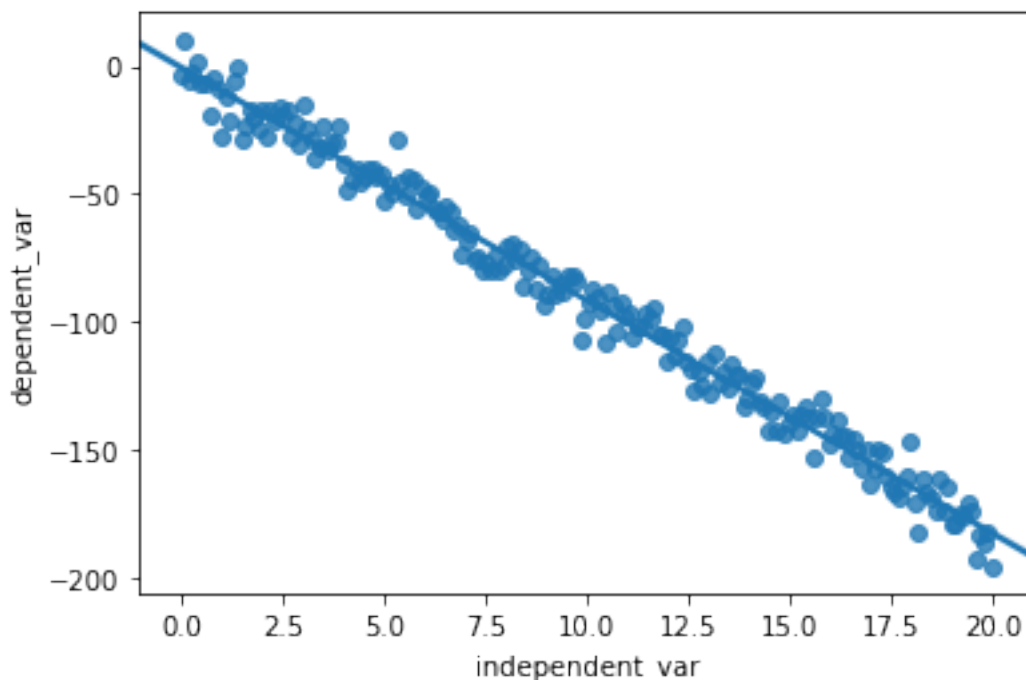
Requirement already satisfied: pytz>=2011k in /usr/lib/python3/dist-packages (from pandas>=0.19->statsmodels==0.10.0rc2) (2018.3)

```
Requirement already satisfied: six in /usr/lib/python3/dist-packages (from
patsy>=0.4.0->statsmodels==0.10.0rc2) (1.11.0)
Installing collected packages: statsmodels
  Found existing installation: statsmodels 0.9.0
    Uninstalling statsmodels-0.9.0:
      Successfully uninstalled statsmodels-0.9.0
Successfully installed statsmodels-0.10.0rc2
```

```
[20]: # import libraries
import numpy as np, pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
from scipy.optimize import minimize
import scipy.stats as stats
import pymc3 as pm3
import numdifftools as ndt
import statsmodels.api as sm
from statsmodels.base.model import GenericLikelihoodModel

# generate data
N = 200
x = np.linspace(0,20,N)
e = np.random.normal(loc = 0.0, scale = 7.0, size = N)
y = -9*x + e
df = pd.DataFrame({'dependent_var':y, 'independent_var':x})
df['constant'] = 1
```

```
[21]: # plot
sns.regplot(df.independent_var, df.dependent_var);
```



```
[4]: df.head()
```

```
[4]:
```

	y	x	constant
0	11.600880	0.000000	1
1	-3.151675	0.202020	1
2	2.477100	0.404040	1
3	-4.506962	0.606061	1
4	-1.766583	0.808081	1

Before we go into the implementation of log likelihood, please take a quick look at the documentation [Scipy norm documentation](#)

From the documentation we can see that given the location and scale, we can find the log probabilities of the corresponding normal distribution density

```
[17]: print(f"The density values for data points are {stats.norm.pdf([1,-1], loc=0, scale=1)}")
print(f"The natural log of the density values for the data points are {stats.norm.logpdf([1,-1], loc=0, scale=1)}")
```

The density values for data points are [0.24197072 0.24197072]

The natural log of the density values for the data points are [-1.41893853 -1.41893853]

[Verify the result yourself](#)

## 0.1 Ordinary Least Squares Fit Using Statsmodels

```
[24]: # split features and target
X = df[['constant', 'independent_var']]
# fit model and summarize
statsmodels_ols_summary = sm.OLS(y,X).fit().summary()
statsmodels_ols_summary
```

```
[24]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.985
Model:                  OLS    Adj. R-squared:           0.985
Method:                 Least Squares    F-statistic:        1.281e+04
Date:                  Tue, 06 Aug 2019    Prob (F-statistic):    6.42e-182
Time:                  11:30:02    Log-Likelihood:       -660.15
No. Observations:      200    AIC:                1324.
Df Residuals:          198    BIC:                1331.
Df Model:               1
Covariance Type:       nonrobust
=====
===
                        coef      std err          t      P>|t|      [0.025
0.975]
-----
---
constant              -0.5107      0.930      -0.549      0.583      -2.344
1.323
independent_var       -9.1022      0.080     -113.197      0.000      -9.261
-8.944
=====
Omnibus:               0.705    Durbin-Watson:       1.747
Prob(Omnibus):         0.703    Jarque-Bera (JB):     0.606
Skew:                 -0.135    Prob(JB):             0.739
Kurtosis:              3.003    Cond. No.             23.2
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
      """
```

## 0.2 Implementing the Likelihood Function

```
[25]: # define likelihood function
def neg_log_likelihood(params):
    """ Computes the """
    intercept, beta, sd = params[0], params[1], params[2] # inputs are guesses
    → at our parameters
    yhat = intercept + beta*x # predictions
    # to find the negative log likelihood, sum the log densities as according
    → to the likelihood function formula
    negative_log_likelihood = -np.sum( stats.norm.logpdf(y, loc=yhat, scale=sd)
    → )
    return(negative_log_likelihood)

# let's start with some random coefficient guesses and optimize
guess = np.array([5,5,2])
results = minimize(neg_log_likelihood, guess, method =
    → 'Nelder-Mead', options={'disp': True})

# drop results into df and round to match statsmodels
resultsdf = pd.DataFrame({'coef':results['x']})
resultsdf.index=['constant','x','sigma']
```

Optimization terminated successfully.

Current function value: 660.147313

Iterations: 186

Function evaluations: 339

```
[19]: # Let's also see if there are other Scipy options available for optimization
print(minimize.__doc__)
```

Minimization of scalar function of one or more variables.

Parameters

-----

fun : callable

The objective function to be minimized.

``fun(x, \*args) -> float``

where x is an 1-D array with shape (n,) and `args`  
is a tuple of the fixed parameters needed to completely  
specify the function.

x0 : ndarray, shape (n,)

Initial guess. Array of real elements of size (n,),  
where 'n' is the number of independent variables.

args : tuple, optional

Extra arguments passed to the objective function and its derivatives (`fun`, `jac` and `hess` functions).

method : str or callable, optional

Type of solver. Should be one of

- 'Nelder-Mead' :ref:`(see here) <optimize.minimize-neldermead>`
- 'Powell' :ref:`(see here) <optimize.minimize-powell>`
- 'CG' :ref:`(see here) <optimize.minimize-cg>`
- 'BFGS' :ref:`(see here) <optimize.minimize-bfgs>`
- 'Newton-CG' :ref:`(see here) <optimize.minimize-newtoncg>`
- 'L-BFGS-B' :ref:`(see here) <optimize.minimize-lbfgsb>`
- 'TNC' :ref:`(see here) <optimize.minimize-tnc>`
- 'COBYLA' :ref:`(see here) <optimize.minimize-cobyla>`
- 'SLSQP' :ref:`(see here) <optimize.minimize-slsqp>`
- 'trust-constr':ref:`(see here) <optimize.minimize-trustconstr>`
- 'dogleg' :ref:`(see here) <optimize.minimize-dogleg>`
- 'trust-ncg' :ref:`(see here) <optimize.minimize-trustncg>`
- 'trust-exact' :ref:`(see here) <optimize.minimize-trustexact>`
- 'trust-krylov' :ref:`(see here) <optimize.minimize-trustkrylov>`
- custom - a callable object (added in version 0.14.0),  
see below for description.

If not given, chosen to be one of ``BFGS``, ``L-BFGS-B``, ``SLSQP``, depending if the problem has constraints or bounds.

jac : {callable, '2-point', '3-point', 'cs', bool}, optional

Method for computing the gradient vector. Only for CG, BFGS, Newton-CG, L-BFGS-B, TNC, SLSQP, dogleg, trust-ncg, trust-krylov, trust-exact and trust-constr. If it is a callable, it should be a function that returns the gradient vector:

```
``jac(x, *args) -> array_like, shape (n,)``
```

where x is an array with shape (n,) and `args` is a tuple with the fixed parameters. Alternatively, the keywords {'2-point', '3-point', 'cs'} select a finite difference scheme for numerical estimation of the gradient. Options '3-point' and 'cs' are available only to 'trust-constr'. If `jac` is a Boolean and is True, `fun` is assumed to return the gradient along with the objective function. If False, the gradient will be estimated using '2-point' finite difference estimation.

hess : {callable, '2-point', '3-point', 'cs', HessianUpdateStrategy}, optional

Method for computing the Hessian matrix. Only for Newton-CG, dogleg, trust-ncg, trust-krylov, trust-exact and trust-constr. If it is callable, it should return the Hessian matrix:

```
``hess(x, *args) -> {LinearOperator, spmatrix, array}, (n, n)``
```

where `x` is a `(n,)` ndarray and ``args`` is a tuple with the fixed parameters. `LinearOperator` and sparse matrix returns are allowed only for 'trust-constr' method. Alternatively, the keywords {'2-point', '3-point', 'cs'} select a finite difference scheme for numerical estimation. Or, objects implementing ``HessianUpdateStrategy`` interface can be used to approximate the Hessian. Available quasi-Newton methods implementing this interface are:

- ``BFGS``;
- ``SR1``.

Whenever the gradient is estimated via finite-differences, the Hessian cannot be estimated with options {'2-point', '3-point', 'cs'} and needs to be estimated using one of the quasi-Newton strategies. Finite-difference options {'2-point', '3-point', 'cs'} and ``HessianUpdateStrategy`` are available only for 'trust-constr' method.

`hessp` : callable, optional  
Hessian of objective function times an arbitrary vector `p`. Only for Newton-CG, trust-ncg, trust-krylov, trust-constr.  
Only one of ``hessp`` or ``hess`` needs to be given. If ``hess`` is provided, then ``hessp`` will be ignored. ``hessp`` must compute the Hessian times an arbitrary vector:

```
``hessp(x, p, *args) -> ndarray shape (n,)``
```

where `x` is a `(n,)` ndarray, `p` is an arbitrary vector with dimension `(n,)` and ``args`` is a tuple with the fixed parameters.

`bounds` : sequence or ``Bounds``, optional  
Bounds on variables for L-BFGS-B, TNC, SLSQP and trust-constr methods. There are two ways to specify the bounds:

1. Instance of ``Bounds`` class.
2. Sequence of ```(min, max)``` pairs for each element in ``x``. None is used to specify no bound.

`constraints` : {Constraint, dict} or List of {Constraint, dict}, optional  
Constraints definition (only for COBYLA, SLSQP and trust-constr).  
Constraints for 'trust-constr' are defined as a single object or a list of objects specifying constraints to the optimization problem. Available constraints are:

- ``LinearConstraint``
- ``NonlinearConstraint``

Constraints for COBYLA, SLSQP are defined as a list of dictionaries.

Each dictionary with fields:

```
type : str
    Constraint type: 'eq' for equality, 'ineq' for inequality.
fun : callable
    The function defining the constraint.
jac : callable, optional
    The Jacobian of `fun` (only for SLSQP).
args : sequence, optional
    Extra arguments to be passed to the function and Jacobian.
```

Equality constraint means that the constraint function result is to be zero whereas inequality means that it is to be non-negative.

Note that COBYLA only supports inequality constraints.

```
tol : float, optional
    Tolerance for termination. For detailed control, use solver-specific options.
```

```
options : dict, optional
    A dictionary of solver options. All methods accept the following generic options:
```

```
maxiter : int
    Maximum number of iterations to perform.
disp : bool
    Set to True to print convergence messages.
```

For method-specific options, see :func:`show\_options`.

```
callback : callable, optional
    Called after each iteration. For 'trust-constr' it is a callable with the signature:
```

```
``callback(xk, OptimizeResult state) -> bool``
```

where ``xk`` is the current parameter vector. and ``state`` is an `OptimizeResult` object, with the same fields as the ones from the return. If callback returns True the algorithm execution is terminated.

For all the other methods, the signature is:

```
``callback(xk)``
```

where ``xk`` is the current parameter vector.

Returns

-----

```
res : OptimizeResult
```

The optimization result represented as a ``OptimizeResult`` object. Important attributes are: ``x`` the solution array, ``success`` a



Boolean flag indicating if the optimizer exited successfully and  
``message`` which describes the cause of the termination. See  
`OptimizeResult` for a description of other attributes.

See also

-----

minimize\_scalar : Interface to minimization algorithms for scalar  
univariate functions  
show\_options : Additional options accepted by the solvers

Notes

-----

This section describes the available solvers that can be selected by the  
'method' parameter. The default method is *\*BFGS\**.

**\*\*Unconstrained minimization\*\***

Method :ref:`Nelder-Mead <optimize.minimize-neldermead>` uses the  
Simplex algorithm [1]\_, [2]\_. This algorithm is robust in many  
applications. However, if numerical computation of derivative can be  
trusted, other algorithms using the first and/or second derivatives  
information might be preferred for their better performance in  
general.

Method :ref:`Powell <optimize.minimize-powell>` is a modification  
of Powell's method [3]\_, [4]\_ which is a conjugate direction  
method. It performs sequential one-dimensional minimizations along  
each vector of the directions set (`direc` field in `options` and  
`info`), which is updated at each iteration of the main  
minimization loop. The function need not be differentiable, and no  
derivatives are taken.

Method :ref:`CG <optimize.minimize-cg>` uses a nonlinear conjugate  
gradient algorithm by Polak and Ribiere, a variant of the  
Fletcher-Reeves method described in [5]\_ pp. 120-122. Only the  
first derivatives are used.

Method :ref:`BFGS <optimize.minimize-bfgs>` uses the quasi-Newton  
method of Broyden, Fletcher, Goldfarb, and Shanno (BFGS) [5]\_  
pp. 136. It uses the first derivatives only. BFGS has proven good  
performance even for non-smooth optimizations. This method also  
returns an approximation of the Hessian inverse, stored as  
`hess\_inv` in the `OptimizeResult` object.

Method :ref:`Newton-CG <optimize.minimize-newtoncg>` uses a  
Newton-CG algorithm [5]\_ pp. 168 (also known as the truncated  
Newton method). It uses a CG method to compute the search  
direction. See also *\*TNC\** method for a box-constrained

minimization with a similar algorithm. Suitable for large-scale problems.

Method :ref:`dogleg <optimize.minimize-dogleg>` uses the dog-leg trust-region algorithm [5]\_ for unconstrained minimization. This algorithm requires the gradient and Hessian; furthermore the Hessian is required to be positive definite.

Method :ref:`trust-ncg <optimize.minimize-trustncg>` uses the Newton conjugate gradient trust-region algorithm [5]\_ for unconstrained minimization. This algorithm requires the gradient and either the Hessian or a function that computes the product of the Hessian with a given vector. Suitable for large-scale problems.

Method :ref:`trust-krylov <optimize.minimize-trustkrylov>` uses the Newton GLTR trust-region algorithm [14]\_, [15]\_ for unconstrained minimization. This algorithm requires the gradient and either the Hessian or a function that computes the product of the Hessian with a given vector. Suitable for large-scale problems. On indefinite problems it requires usually less iterations than the `trust-ncg` method and is recommended for medium and large-scale problems.

Method :ref:`trust-exact <optimize.minimize-trustexact>` is a trust-region method for unconstrained minimization in which quadratic subproblems are solved almost exactly [13]\_. This algorithm requires the gradient and the Hessian (which is *\*not\** required to be positive definite). It is, in many situations, the Newton method to converge in fewer iteration and the most recommended for small and medium-size problems.

#### **\*\*Bound-Constrained minimization\*\***

Method :ref:`L-BFGS-B <optimize.minimize-lbfgsb>` uses the L-BFGS-B algorithm [6]\_, [7]\_ for bound constrained minimization.

Method :ref:`TNC <optimize.minimize-tnc>` uses a truncated Newton algorithm [5]\_, [8]\_ to minimize a function with variables subject to bounds. This algorithm uses gradient information; it is also called Newton Conjugate-Gradient. It differs from the *\*Newton-CG\** method described above as it wraps a C implementation and allows each variable to be given upper and lower bounds.

#### **\*\*Constrained Minimization\*\***

Method :ref:`COBYLA <optimize.minimize-cobyla>` uses the Constrained Optimization BY Linear Approximation (COBYLA) method [9]\_, [10]\_, [11]\_. The algorithm is based on linear approximations to the objective function and each constraint. The

method wraps a FORTRAN implementation of the algorithm. The constraints functions 'fun' may return either a single number or an array or list of numbers.

Method :ref:`SLSQP <optimize.minimize-slsqp>` uses Sequential Least Squares Programming to minimize a function of several variables with any combination of bounds, equality and inequality constraints. The method wraps the SLSQP Optimization subroutine originally implemented by Dieter Kraft [12]\_. Note that the wrapper handles infinite values in bounds by converting them into large floating values.

Method :ref:`trust-constr <optimize.minimize-trustconstr>` is a trust-region algorithm for constrained optimization. It switches between two implementations depending on the problem definition. It is the most versatile constrained minimization algorithm implemented in SciPy and the most appropriate for large-scale problems. For equality constrained problems it is an implementation of Byrd-Omojokun Trust-Region SQP method described in [17]\_ and in [5]\_, p. 549. When inequality constraints are imposed as well, it switches to the trust-region interior point method described in [16]\_. This interior point algorithm, in turn, solves inequality constraints by introducing slack variables and solving a sequence of equality-constrained barrier problems for progressively smaller values of the barrier parameter. The previously described equality constrained SQP method is used to solve the subproblems with increasing levels of accuracy as the iterate gets closer to a solution.

#### **\*\*Finite-Difference Options\*\***

For Method :ref:`trust-constr <optimize.minimize-trustconstr>` the gradient and the Hessian may be approximated using three finite-difference schemes: {'2-point', '3-point', 'cs'}. The scheme 'cs' is, potentially, the most accurate but it requires the function to correctly handles complex inputs and to be differentiable in the complex plane. The scheme '3-point' is more accurate than '2-point' but requires twice as much operations.

#### **\*\*Custom minimizers\*\***

It may be useful to pass a custom minimization method, for example when using a frontend to this method such as ``scipy.optimize.basinhopping`` or a different library. You can simply pass a callable as the ``method`` parameter.

The callable is called as ``method(fun, x0, args, **kwargs, **options)`` where ``kwargs`` corresponds to any other parameters passed to ``minimize`` (such as ``callback``, ``hess``, etc.), except the ``options`` dict, which has

its contents also passed as ``method`` parameters pair by pair. Also, if ``jac`` has been passed as a bool type, ``jac`` and ``fun`` are mangled so that ``fun`` returns just the function values and ``jac`` is converted to a function returning the Jacobian. The method shall return an ``OptimizeResult`` object.

The provided ``method`` callable must be able to accept (and possibly ignore) arbitrary parameters; the set of parameters accepted by ``minimize`` may expand in future versions and then these parameters will be passed to the method. You can find an example in the `scipy.optimize` tutorial.

.. versionadded:: 0.11.0

## References

-----

- .. [1] Nelder, J A, and R Mead. 1965. A Simplex Method for Function Minimization. The Computer Journal 7: 308-13.
- .. [2] Wright M H. 1996. Direct search methods: Once scorned, now respectable, in Numerical Analysis 1995: Proceedings of the 1995 Dundee Biennial Conference in Numerical Analysis (Eds. D F Griffiths and G A Watson). Addison Wesley Longman, Harlow, UK. 191-208.
- .. [3] Powell, M J D. 1964. An efficient method for finding the minimum of a function of several variables without calculating derivatives. The Computer Journal 7: 155-162.
- .. [4] Press W, S A Teukolsky, W T Vetterling and B P Flannery. Numerical Recipes (any edition), Cambridge University Press.
- .. [5] Nocedal, J, and S J Wright. 2006. Numerical Optimization. Springer New York.
- .. [6] Byrd, R H and P Lu and J. Nocedal. 1995. A Limited Memory Algorithm for Bound Constrained Optimization. SIAM Journal on Scientific and Statistical Computing 16 (5): 1190-1208.
- .. [7] Zhu, C and R H Byrd and J Nocedal. 1997. L-BFGS-B: Algorithm 778: L-BFGS-B, FORTRAN routines for large scale bound constrained optimization. ACM Transactions on Mathematical Software 23 (4): 550-560.
- .. [8] Nash, S G. Newton-Type Minimization Via the Lanczos Method. 1984. SIAM Journal of Numerical Analysis 21: 770-778.
- .. [9] Powell, M J D. A direct search optimization method that models the objective and constraint functions by linear interpolation. 1994. Advances in Optimization and Numerical Analysis, eds. S. Gomez and J-P Hennart, Kluwer Academic (Dordrecht), 51-67.
- .. [10] Powell M J D. Direct search algorithms for optimization calculations. 1998. Acta Numerica 7: 287-336.
- .. [11] Powell M J D. A view of algorithms for optimization without derivatives. 2007. Cambridge University Technical Report DAMTP 2007/NA03
- .. [12] Kraft, D. A software package for sequential quadratic

- programming. 1988. Tech. Rep. DFVLR-FB 88-28, DLR German Aerospace Center -- Institute for Flight Mechanics, Koln, Germany.
- .. [13] Conn, A. R., Gould, N. I., and Toint, P. L.  
Trust region methods. 2000. Siam. pp. 169-200.
  - .. [14] F. Lenders, C. Kirches, A. Potschka: "trlib: A vector-free implementation of the GLTR method for iterative solution of the trust region problem", <https://arxiv.org/abs/1611.04718>
  - .. [15] N. Gould, S. Lucidi, M. Roma, P. Toint: "Solving the Trust-Region Subproblem using the Lanczos Method", SIAM J. Optim., 9(2), 504--525, (1999).
  - .. [16] Byrd, Richard H., Mary E. Hribar, and Jorge Nocedal. 1999.  
An interior point algorithm for large-scale nonlinear programming. SIAM Journal on Optimization 9.4: 877-900.
  - .. [17] Lalee, Marucha, Jorge Nocedal, and Todd Plantega. 1998. On the implementation of an algorithm for large-scale equality constrained optimization. SIAM Journal on Optimization 8.3: 682-706.

## Examples

-----

Let us consider the problem of minimizing the Rosenbrock function. This function (and its respective derivatives) is implemented in ``rosen`` (resp. ``rosen_der``, ``rosen_hess``) in the ``scipy.optimize``.

```
>>> from scipy.optimize import minimize, rosen, rosen_der
```

A simple application of the *\*Nelder-Mead\** method is:

```
>>> x0 = [1.3, 0.7, 0.8, 1.9, 1.2]
>>> res = minimize(rosen, x0, method='Nelder-Mead', tol=1e-6)
>>> res.x
array([ 1.,  1.,  1.,  1.,  1.])
```

Now using the *\*BFGS\** algorithm, using the first derivative and a few options:

```
>>> res = minimize(rosen, x0, method='BFGS', jac=rosen_der,
...               options={'gtol': 1e-6, 'disp': True})
Optimization terminated successfully.
    Current function value: 0.000000
    Iterations: 26
    Function evaluations: 31
    Gradient evaluations: 31
>>> res.x
array([ 1.,  1.,  1.,  1.,  1.])
>>> print(res.message)
Optimization terminated successfully.
>>> res.hess_inv
array([[ 0.00749589,  0.01255155,  0.02396251,  0.04750988,  0.09495377], #
```

may vary

```
[ 0.01255155,  0.02510441,  0.04794055,  0.09502834,  0.18996269],
[ 0.02396251,  0.04794055,  0.09631614,  0.19092151,  0.38165151],
[ 0.04750988,  0.09502834,  0.19092151,  0.38341252,  0.7664427 ],
[ 0.09495377,  0.18996269,  0.38165151,  0.7664427,   1.53713523]])
```

Next, consider a minimization problem with several constraints (namely Example 16.4 from [5]). The objective function is:

```
>>> fun = lambda x: (x[0] - 1)**2 + (x[1] - 2.5)**2
```

There are three constraints defined as:

```
>>> cons = ({'type': 'ineq', 'fun': lambda x: x[0] - 2 * x[1] + 2},
...         {'type': 'ineq', 'fun': lambda x: -x[0] - 2 * x[1] + 6},
...         {'type': 'ineq', 'fun': lambda x: -x[0] + 2 * x[1] + 2})
```

And variables must be positive, hence the following bounds:

```
>>> bnds = ((0, None), (0, None))
```

The optimization problem is solved using the SLSQP method as:

```
>>> res = minimize(fun, (2, 0), method='SLSQP', bounds=bnds,
...                constraints=cons)
```

It should converge to the theoretical solution (1.4 ,1.7).

[10]: statsmodels\_ols\_summary

[10]: <class 'statsmodels.iolib.summary.Summary'>  
"""

```

                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.982
Model:                  OLS      Adj. R-squared:          0.982
Method:                 Least Squares      F-statistic:        1.080e+04
Date:                   Tue, 06 Aug 2019      Prob (F-statistic):    1.13e-174
Time:                   11:13:41      Log-Likelihood:       -672.67
No. Observations:       200      AIC:                  1349.
Df Residuals:           198      BIC:                  1356.
Df Model:                1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
=====
```

```

-----
constant      -1.4593      0.990      -1.474      0.142      -3.411      0.493
x              -8.8947      0.086     -103.903      0.000      -9.064     -8.726
=====
Omnibus:                0.634   Durbin-Watson:                1.803
Prob(Omnibus):          0.728   Jarque-Bera (JB):          0.415
Skew:                   -0.099   Prob(JB):                  0.813
Kurtosis:               3.101   Cond. No.                  23.2
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""

```

```
[26]: np.round(resultsdf.head(2), 4)
```

```

[26]:          coef
constant -0.5108
x        -9.1022

```

```
[ ]:
```